

Getting Started with OpenACC

Justin Luitjens, NVIDIA

3 Ways to Accelerate Applications

Applications

Libraries

“Drop-in”
Acceleration

OpenACC
Directives

Easily Accelerate
Applications

Programming
Languages

Maximum
Flexibility

OpenACC

The Standard for GPU Directives

- **Simple:** Directives are the easy path to accelerate compute intensive applications
- **Open:** OpenACC is an open GPU directives standard, making GPU programming straightforward and portable across parallel and multi-core processors
- **Powerful:** GPU Directives allow complete access to the massive parallel power of a GPU

Directives: Easy & Powerful



Real-Time Object Detection

Global Manufacturer of Navigation Systems



5x in 40 Hours

Valuation of Stock Portfolios using Monte Carlo


Global Technology Consulting Company



2x in 4 Hours

Interaction of Solvents and Biomolecules

University of Texas at San Antonio



5x in 8 Hours

“Optimizing code with directives is quite easy, especially compared to CPU threads or writing CUDA kernels. The most important thing is avoiding restructuring of existing code for production applications.”

Focus on Expressing Parallelism




With Directives, tuning work focuses on *expressing parallelism*, which makes codes inherently better


Example: Application tuning work using directives for new Titan system at ORNL

S3D

Research more efficient combustion with next-generation fuels



- Tuning top 3 kernels (90% of runtime)
- 3 to 6x faster on CPU+GPU vs. CPU+CPU
- But also improved all-CPU version by 50%



CAM-SE

Answer questions about specific climate change adaptation and mitigation scenarios

- Tuning top key kernel (50% of runtime)
- 6.5x faster on CPU+GPU vs. CPU+CPU
- Improved performance of CPU version by 100%

OpenACC is not
GPU Programming.

OpenACC is
Expressing Parallelism
in your code.

OpenACC Specification and Website



- Full OpenACC 1.0 Specification available online

www.openacc.org

- Quick reference card also available
- Compilers available now from PGI, Cray, and CAPS

The OpenACC™ API QUICK REFERENCE GUIDE

The OpenACC Application Program Interface describes a collection of compiler directives to specify loops and regions of code in standard C, C++ and Fortran to be offloaded from a host CPU to an attached accelerator, providing portability across operating systems, host CPUs and accelerators.

Most OpenACC directives apply to the immediately following structured block or loop; a structured block is a single statement or a compound statement (C or C++) or a sequence of statements (Fortran) with a single entry point at the top and a single exit at the bottom.



PGI

Version 1.0, November 2011

Start Now with OpenACC Directives



Sign up for a **free trial** of the directives compiler now!

Free trial license to PGI Accelerator

Tools for quick ramp

www.nvidia.com/gpudirectives



GPU COMPUTING SOLUTIONS

- Main
- What is GPU Computing?
- Why Choose Tesla
- Industry Software Solutions
- Tesla Workstation Solutions
- Tesla Data Center Solutions
- Tesla Bio Workbench
- Where to Buy
- Contact US
- Sign up for Tesla Alerts
- Fermi GPU Computing Architecture

SOFTWARE AND HARDWARE INFO

- Tesla Product Literature
- Tesla Software Features
- Software Development Tools
- CUDA Training and Consulting Services
- GPU Cloud Computing Service Providers
- OpenACC GPU Directives

Accelerate Your Scientific Code with OpenACC

The Open Standard for GPU Accelerator Directives

Thousands of cores working for you.

Based on the [OpenACC](#) standard, GPU directives are the easy, proven way to accelerate your scientific or industrial code. With GPU directives, you can accelerate your code by simply inserting compiler hints into your code and the compiler will automatically map compute-intensive portions of your code to the GPU. Here's an example of how easy a single directive hint can accelerate the calculation of pi. With GPU directives, you can get started and see results in the same afternoon.

```
#include <stdio.h>
#define N 10000
int main(void) {
    double pi = 0.0f; long i;
    #pragma acc region for
    for (i=0; i<N; i++)
    {
        double t= (double) ((i+0.5)/N);
        pi +=4.0/(1.0+t*t);
    }
    printf("pi=%f\n",pi/N);
    return 0;
}
```

By starting with a free, 30-day trial of PGI directives today, you are working on the technology that is the foundation of the OpenACC directives standard. OpenACC is:

"I have written micron (written in Fortran 90) properties of two and dimensional magnetic directives approach error my existing code perform my computation which resulted in a speedup (more than 20 computation." [Learn more](#)

Professor M. Amin Kay
University of Houston

"The PGI compiler is not just how powerful it is software we are writing times faster on the NV are very pleased and future uses. It's like on supercomputer." [Learn more](#)

Dr. Kerry Black
University of Melbourne

Expressing Parallelism with OpenACC

A Very Simple Exercise: SAXPY



SAXPY in C

```
void saxpy(int n,
           float a,
           float *x,
           float *restrict y)
{
    for (int i = 0; i < n; ++i)
        y[i] = a*x[i] + y[i];
}

...
// Perform SAXPY on 1M elements
saxpy(1<<20, 2.0, x, y);
...
```

SAXPY in Fortran

```
subroutine saxpy(n, a, x, y)
    real :: x(n), y(n), a
    integer :: n, i

    do i=1,n
        y(i) = a*x(i)+y(i)
    enddo

end subroutine saxpy

...
! Perform SAXPY on 1M elements
call saxpy(2**20, 2.0, x_d,
y_d)
...
```

A Very Simple Exercise: SAXPY OpenMP



SAXPY in C

```
void saxpy(int n,
           float a,
           float *x,
           float *restrict y)
{
    #pragma omp parallel for
    for (int i = 0; i < n; ++i)
        y[i] = a*x[i] + y[i];
}

...
// Perform SAXPY on 1M elements
saxpy(1<<20, 2.0, x, y);
...
```

SAXPY in Fortran

```
subroutine saxpy(n, a, x, y)
    real :: x(n), y(n), a
    integer :: n, i

    !$omp parallel do
    do i=1,n
        y(i) = a*x(i)+y(i)
    enddo
    !$omp end parallel do
end subroutine saxpy

...
! Perform SAXPY on 1M elements
call saxpy(2**20, 2.0, x_d,
y_d)
...
```

A Very Simple Exercise: SAXPY OpenACC



SAXPY in C

```
void saxpy(int n,
           float a,
           float *x,
           float *restrict y)
{
    #pragma acc parallel loop
    for (int i = 0; i < n; ++i)
        y[i] = a*x[i] + y[i];
}

...
// Perform SAXPY on 1M elements
saxpy(1<<20, 2.0, x, y);
...
```

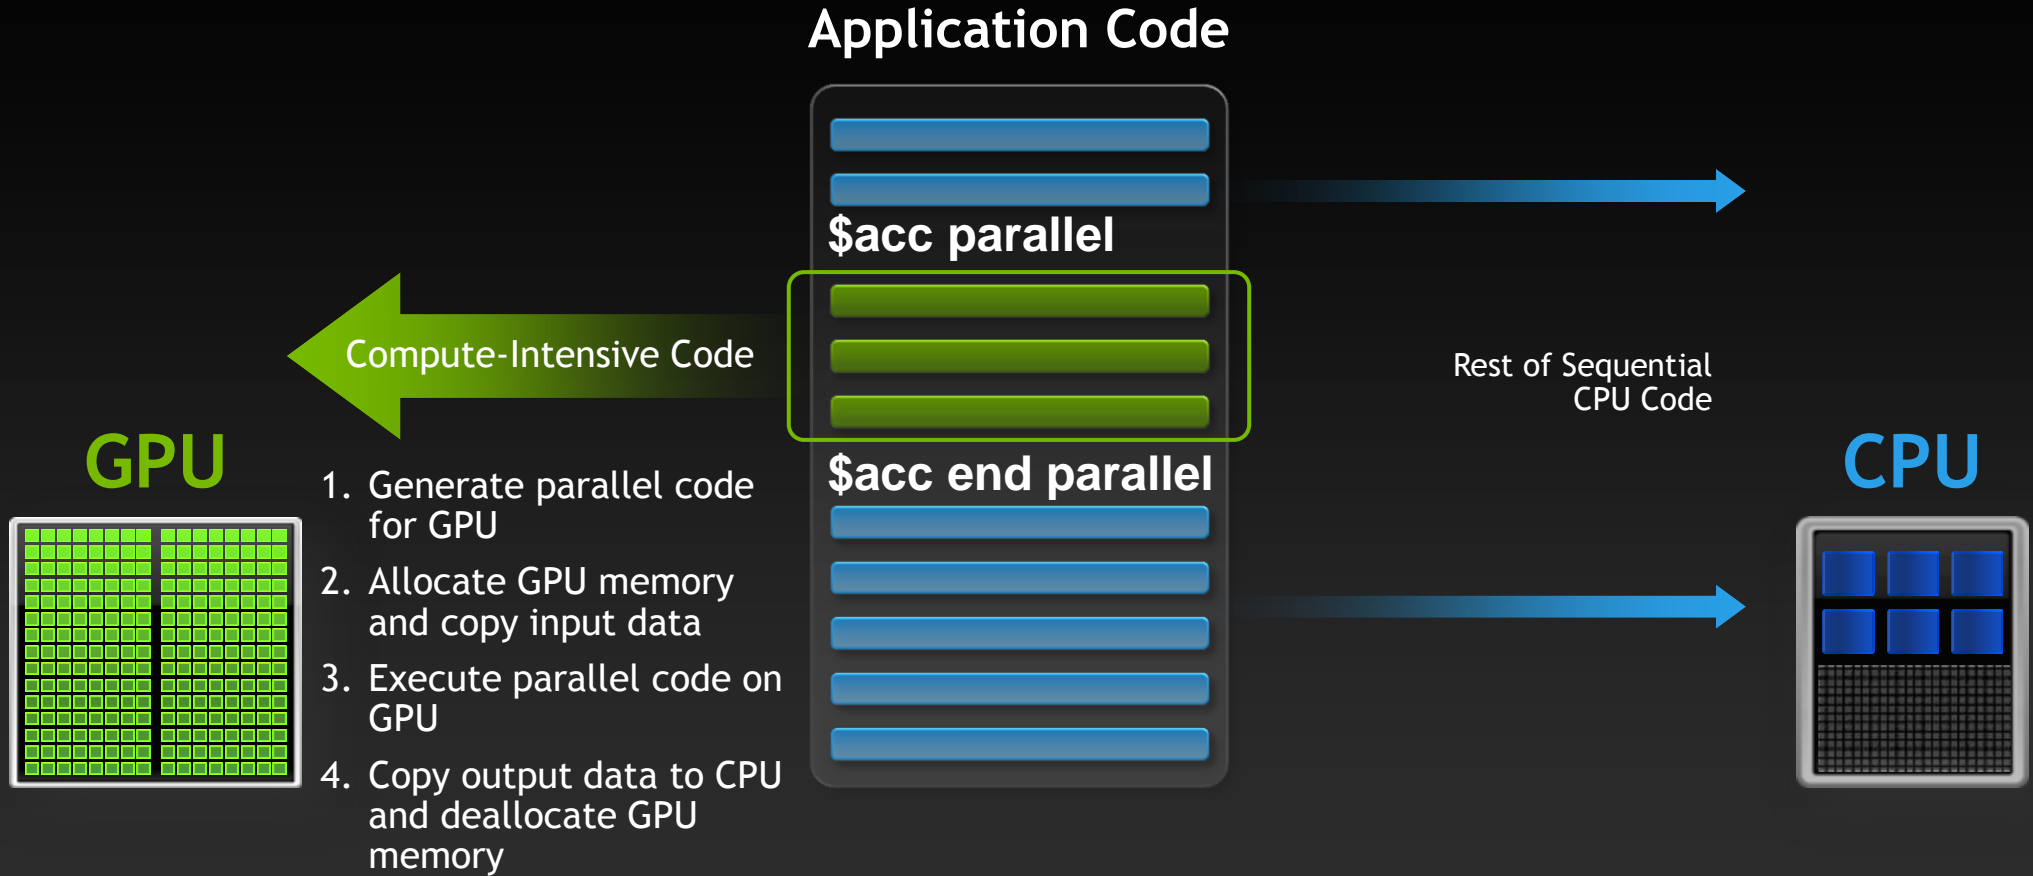
SAXPY in Fortran

```
subroutine saxpy(n, a, x, y)
    real :: x(n), y(n), a
    integer :: n, i

    !$acc parallel loop
    do i=1,n
        y(i) = a*x(i)+y(i)
    enddo
    !$acc end parallel loop
end subroutine saxpy

...
! Perform SAXPY on 1M elements
call saxpy(2**20, 2.0, x_d,
y_d)
...
```

OpenACC Execution Model



Directive Syntax



- Fortran

!\$acc directive [clause [,] clause] ...]

...often paired with a matching end directive surrounding a structured code block:

!\$acc end directive

- C

#pragma acc directive [clause [,] clause] ...]

...often followed by a structured code block

- Common Clauses

if(condition) , async(handle)

Complete SAXPY example code



- Trivial first example
 - Apply a loop directive
 - Learn compiler commands

```
#include <stdlib.h>

void saxpy(int n,
           float a,
           float *x,
           float *restrict y)
{
    #pragma acc parallel loop
    for (int i = 0; i < n; ++i)
        y[i] = a * x[i] + y[i];
}
```

```
int main(int argc, char **argv)
{
    int N = 1<<20; // 1 million floats

    if (argc > 1)
        N = atoi(argv[1]);

    float *x = (float*)malloc(N * sizeof(float));
    float *y = (float*)malloc(N * sizeof(float));

    for (int i = 0; i < N; ++i) {
        x[i] = 2.0f;
        y[i] = 1.0f;
    }

    saxpy(N, 3.0f, x, y);

    return 0;
}
```

Compile (PGI)



- C:

```
pgcc -acc [-Minfo=accel] [-ta=nvidia] -o saxpy_acc saxpy.c
```

- Fortran:

```
pgf90 -acc [-Minfo=accel] [-ta=nvidia] -o saxpy_acc saxpy.f90
```

- Compiler output:

```
pgcc -acc -Minfo=accel -ta=nvidia -o saxpy_acc saxpy.c
saxpy:
  11, Accelerator kernel generated
    13, #pragma acc loop gang, vector(256) /* blockIdx.x threadIdx.x */
  11, Generating present_or_copyin(x[0:n])
    Generating present_or_copy(y[0:n])
    Generating NVIDIA code
    Generating compute capability 1.0 binary
    Generating compute capability 2.0 binary
    Generating compute capability 3.0 binary
```

- The PGI compiler provides automatic instrumentation when **PGI_ACC_TIME=1** at runtime

```
Accelerator Kernel Timing data
/home/jlarkin/kernels/saxpy/saxpy.c
saxpy NVIDIA devicenum=0
time(us): 3,256
11: data copyin reached 2 times
    device time(us): total=1,619 max=892 min=727 avg=809
11: kernel launched 1 times
    grid: [4096] block: [256]
    device time(us): total=714 max=714 min=714 avg=714
    elapsed time(us): total=724 max=724 min=724 avg=724
15: data copyout reached 1 times
    device time(us): total=923 max=923 min=923 avg=923
```

Another approach: **kernels** construct



- The kernels construct expresses that a region may contain parallelism and the compiler determines what can safely be parallelized.

!\$acc kernels

```
do i=1,n
  a(i) = 0.0
  b(i) = 1.0
  c(i) = 2.0
end do
```

} kernel 1

```
do i=1,n
  a(i) = b(i) + c(i)
end do
```

} kernel 2

!\$acc end kernels

The compiler identifies
2 parallel loops and
generates 2 kernels.

OpenACC parallel vs. kernels



PARALLEL

- Requires analysis by programmer to ensure safe parallelism
- Straightforward path from OpenMP

KERNELS

- Compiler performs parallel analysis and parallelizes what it believes safe
- Can cover larger area of code with single directive


Both approaches are **equally valid** and can perform **equally well**.

OpenACC by Example

Example: Jacobi Iteration



- Iteratively converges to correct value (e.g. Temperature), by computing new values at each point from the average of neighboring points.
 - Common, useful algorithm
 - Example: Solve Laplace equation in 2D: $\nabla^2 f(x, y) = 0$



$$A_{k+1}(i, j) = \frac{A_k(i-1, j) + A_k(i+1, j) + A_k(i, j-1) + A_k(i, j+1)}{4}$$

Jacobi Iteration: C Code



```
while ( err > tol && iter < iter_max ) {  
    err=0.0;  
  
    for( int j = 1; j < n-1; j++) {  
        for(int i = 1; i < m-1; i++) {  
  
            Anew[j][i] = 0.25 * (A[j][i+1] + A[j][i-1] +  
                                A[j-1][i] + A[j+1][i]);  
  
            err = max(err, abs(Anew[j][i] - A[j][i]));  
        }  
    }  
  
    for( int j = 1; j < n-1; j++) {  
        for( int i = 1; i < m-1; i++ ) {  
            A[j][i] = Anew[j][i];  
        }  
    }  
  
    iter++;  
}
```



Iterate until converged



Iterate across matrix
elements



Calculate new value from
neighbors



Compute max error for
convergence



Swap input/output arrays

Jacobi Iteration: OpenMP C Code



```
while ( err > tol && iter < iter_max ) {
    err=0.0;

    #pragma omp parallel for shared(m, n, Anew, A) reduction(max:err)
    for( int j = 1; j < n-1; j++) {
        for(int i = 1; i < m-1; i++) {

            Anew[j][i] = 0.25 * (A[j][i+1] + A[j][i-1] +
                                A[j-1][i] + A[j+1][i]);

            err = max(err, abs(Anew[j][i] - A[j][i]));
        }
    }

    #pragma omp parallel for shared(m, n, Anew, A)
    for( int j = 1; j < n-1; j++) {
        for( int i = 1; i < m-1; i++ ) {
            A[j][i] = Anew[j][i];
        }
    }

    iter++;
}
```

Parallelize loop across
CPU threads

Parallelize loop across
CPU threads

Jacobi Iteration: OpenACC C Code



```
while ( err > tol && iter < iter_max ) {
    err=0.0;

    #pragma acc parallel loop reduction(max:err)
    for( int j = 1; j < n-1; j++) {
        for(int i = 1; i < m-1; i++) {

            Anew[j][i] = 0.25 * (A[j][i+1] + A[j][i-1] +
                                A[j-1][i] + A[j+1][i]);

            err = max(err, abs(Anew[j][i] - A[j][i]));
        }
    }

    #pragma acc parallel loop
    for( int j = 1; j < n-1; j++) {
        for( int i = 1; i < m-1; i++ ) {
            A[j][i] = Anew[j][i];
        }
    }

    iter++;
}
```



Parallelize loop nest on
GPU



Parallelize loop nest on
GPU

PGI Accelerator Compiler output (C)



```
pgcc -Minfo=all -ta=nvidia:5.0,cc3x -acc -Minfo=accel -o laplace2d_acc laplace2d.c
main:
```

```
56, Accelerator kernel generated
```

```
57, #pragma acc loop gang /* blockIdx.x */
```

```
59, #pragma acc loop vector(256) /* threadIdx.x */
```

```
56, Generating present_or_copyin(A[0:][0:])
```

```
Generating present_or_copyout(Anew[1:4094][1:4094])
```

```
Generating NVIDIA code
```

```
Generating compute capability 3.0 binary
```

```
59, Loop is parallelizable
```

```
68, Accelerator kernel generated
```

```
69, #pragma acc loop gang /* blockIdx.x */
```

```
71, #pragma acc loop vector(256) /* threadIdx.x */
```

```
68, Generating present_or_copyout(A[1:4094][1:4094])
```

```
Generating present_or_copyin(Anew[1:4094][1:4094])
```

```
Generating NVIDIA code
```

```
Generating compute capability 3.0 binary
```


```
71, Loop is parallelizable
```

Execution Time (lower is better)



CPU: Intel i7-3930K
6 Cores @ 3.20GHz

GPU: NVIDIA Tesla K20



What went wrong?

- Set **PGI_ACC_TIME** environment variable to '1'

Accelerator Kernel Timing data

/home/jlarkin/openacc-workshop/exercises/001-laplace2D-kernels/laplace2d.c

main NVIDIA devicenum=0

time(us): 93,201,190

56: data copyin reached 1000 times

device time(us): total=23,049,452 max=28,928 min=22,761 avg=23,049

56: kernel launched 1000 times

grid: [4094] block: [256]

device time(us): total=2,609,928 max=2,812 min=2,593

elapsed time(us): total=2,872,585 max=3,022 min=2,642

56: reduction kernel launched 1000 times

grid: [1] block: [256]

device time(us): total=19,218 max=724 min=16 avg=19

elapsed time(us): total=29,070 max=734 min=26 avg=29

68: data copyin reached 1000 times

device time(us): total=23,888,588 max=33,546 min=23,378 avg=23,888

68: kernel launched 1000 times

grid: [4094] block: [256]

device time(us): total=2,398,101 max=2,961 min=2,137 avg=2,398

elapsed time(us): total=2,407,481 max=2,971 min=2,146 avg=2,407

68: data copyout reached 1000 times

device time(us): total=20,664,362 max=27,788 min=20,511 avg=20,664

77: data copyout reached 1000 times

device time(us): total=20,571,541 max=24,837 min=20,521 avg=20,571

23 seconds

2.6 seconds

Huge Data Transfer Bottleneck!
Computation: 5.19 seconds
Data movement: 74.7 seconds

0.19 seconds


23.9 seconds

2.4 seconds

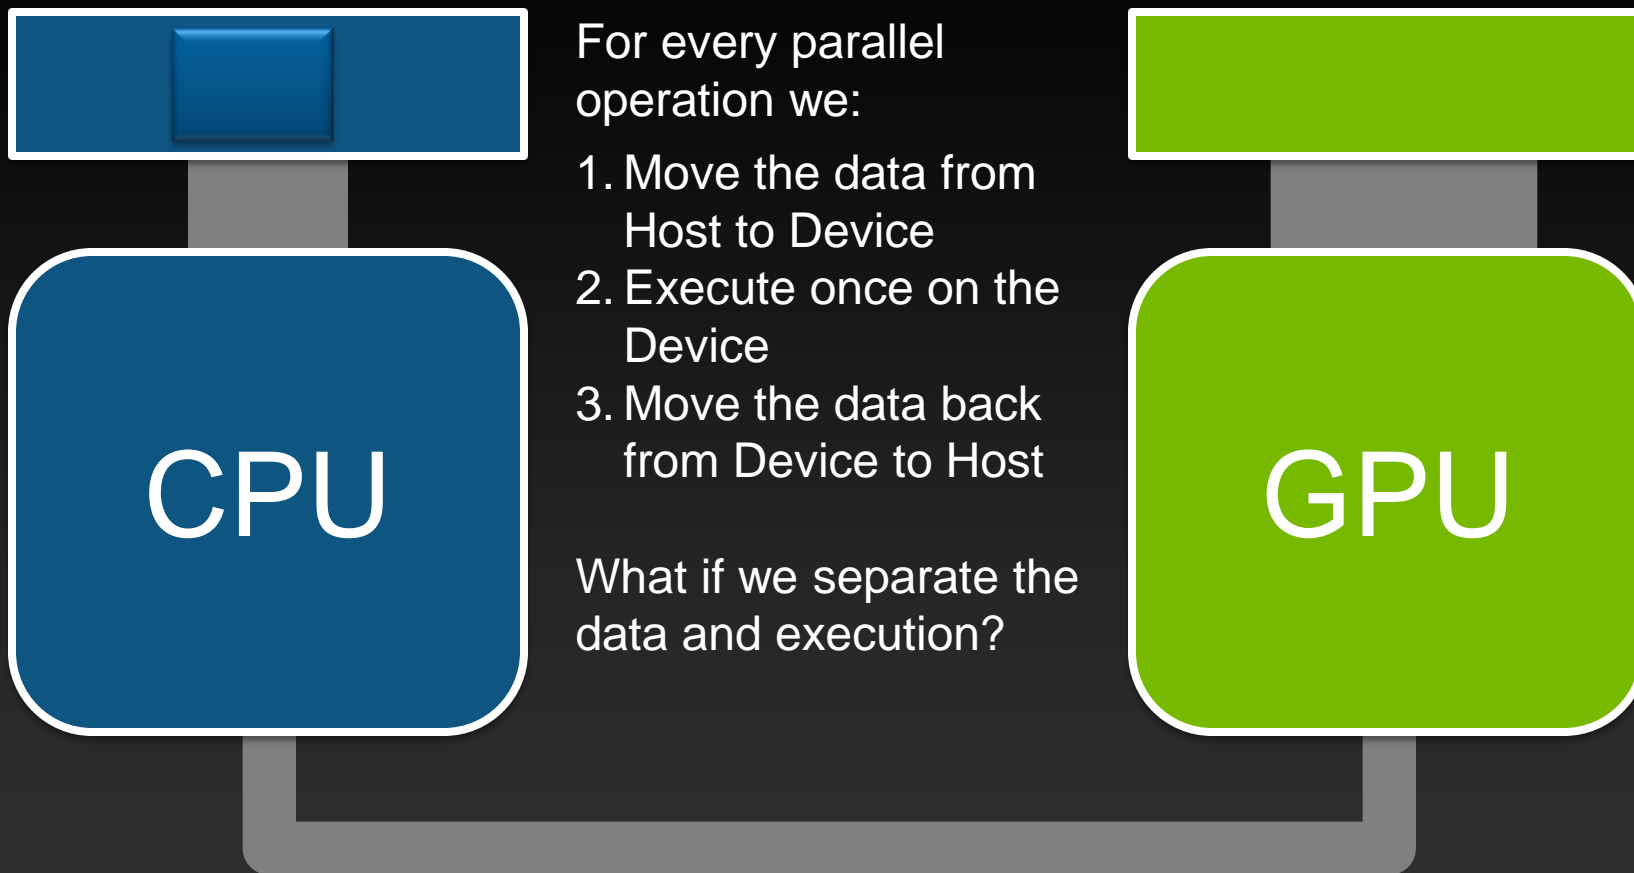
27.8 seconds

24.8 seconds

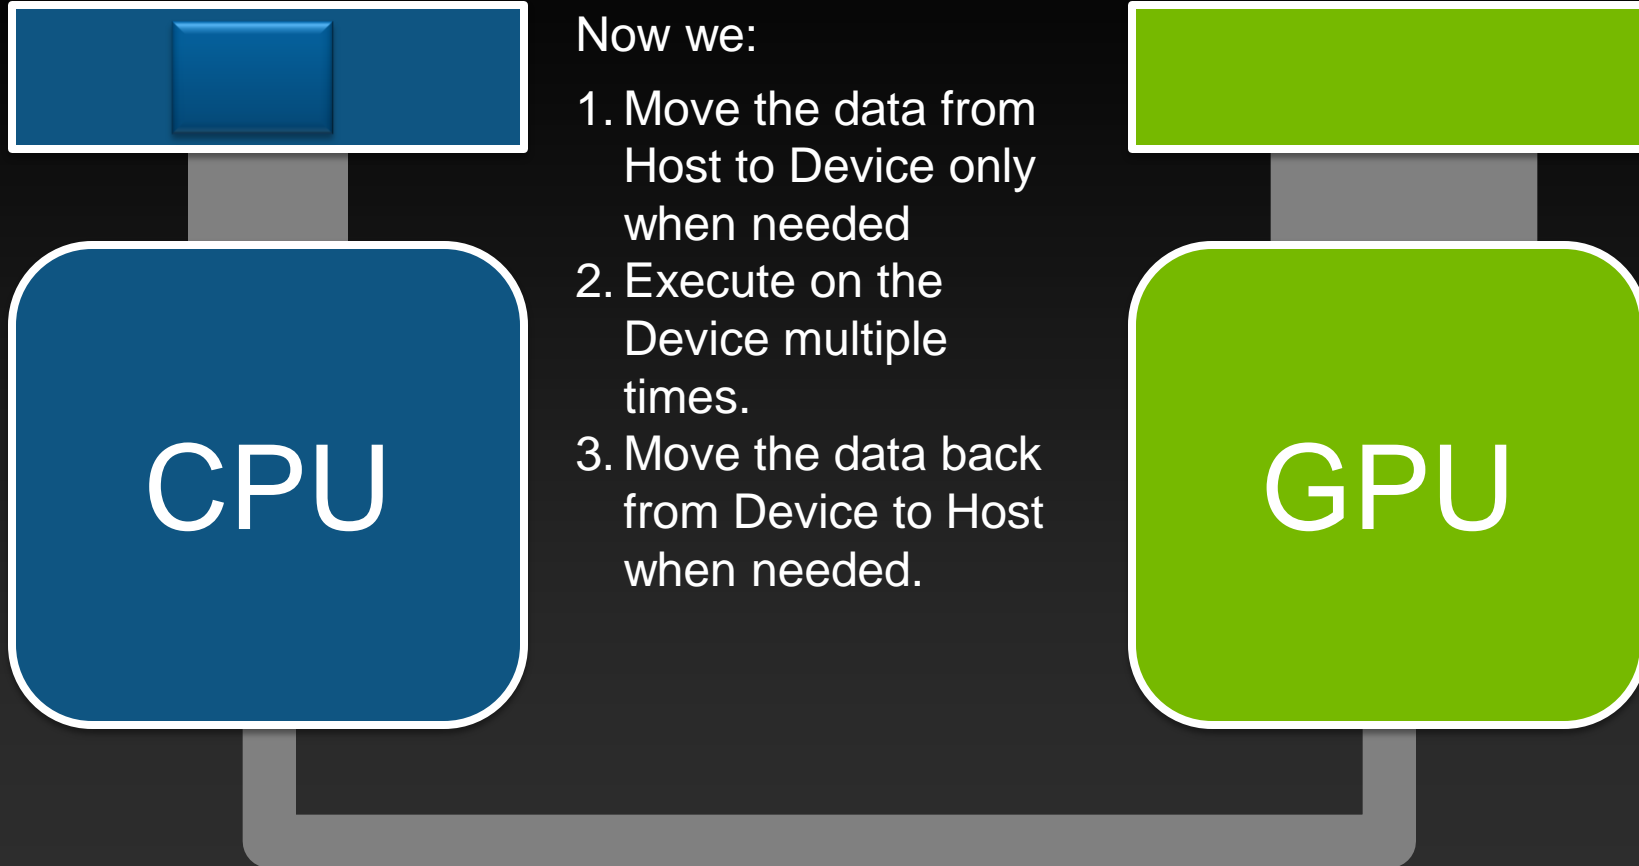
Offloading a Parallel Kernel



Offloading a Parallel Kernel



Separating Data from Computation



Excessive Data Transfers



```
while ( err > tol && iter < iter_max ) {  
    err=0.0;
```

A, Anew resident on host

Copy

```
#pragma acc parallel loop reduction(max:err)
```

A, Anew resident on accelerator

These copies happen
every iteration of the
outer while loop!*

```
for( int j = 1; j < n-1; j++) {  
    for(int i = 1; i < m-1; i++) {  
        Anew[j][i] = 0.25 * (A[j][i+1] + A[j][i-1] +  
                             A[j-1][i] + A[j+1][i]);  
        err = max(err, abs(Anew[j][i] - A[j][i]));  
    }  
}
```

A, Anew resident on host

Copy

A, Anew resident on accelerator

...

```
}
```

And note that there are two `#pragma acc parallel`, so there are 4 copies per while loop iteration!

Data Management with OpenACC

Defining data regions



- The **data** construct defines a region of code in which GPU arrays remain on the GPU and are shared among all kernels in that region.

```
!$acc data
  !$acc parallel loop
  ...
  !$acc parallel loop
  ...
!$acc end data
```

Data Region

Arrays used within the data region will remain on the GPU until the end of the data region.

Data Clauses



- `copy (list)` Allocates memory on GPU and copies data from host to GPU when entering region and copies data to the host when exiting region.
- `copyin (list)` Allocates memory on GPU and copies data from host to GPU when entering region.
- `copyout (list)` Allocates memory on GPU and copies data to the host when exiting region.
- `create (list)` Allocates memory on GPU but does not copy.
- `present (list)` Data is already present on GPU from another containing data region.
- `and present_or_copy[in|out], present_or_create, deviceptr.`

Array Shaping



- Compiler sometimes cannot determine size of arrays
 - Must specify explicitly using data clauses and array “shape”
- C

```
#pragma acc data copyin(a[0:size]), copyout(b[s/4:3*s/4])
```
- Fortran

```
!$acc data copyin(a(1:end)), copyout(b(s/4:3*s/4))
```
- Note: data clauses can be used on data, parallel, or kernels

Jacobi Iteration: Data Directives



- Task: use **acc data** to minimize transfers in the Jacobi example

Jacobi Iteration: OpenACC C Code



Copy A in at beginning of loop, out at end. Allocate Anew on accelerator

```
#pragma acc data copy(A), create(Anew)
while ( err > tol && iter < iter_max ) {
    err=0.0;

    #pragma acc parallel loop reduction(max:err)
    for( int j = 1; j < n-1; j++) {
        for(int i = 1; i < m-1; i++) {

            Anew[j][i] = 0.25 * (A[j][i+1] + A[j][i-1] +
                                A[j-1][i] + A[j+1][i]);

            err = max(err, abs(Anew[j][i] - A[j][i]));
        }
    }

    #pragma acc parallel loop
    for( int j = 1; j < n-1; j++) {
        for( int i = 1; i < m-1; i++ ) {
            A[j][i] = Anew[j][i];
        }
    }

    iter++;
}
```

Did it help?

- Set **PGI_ACC_TIME** environment variable to '1'

Accelerator Kernel Timing data

/home/jlarkin/openacc-workshop/exercises/001-laplace2D-kernels/laplace2d.c

main NVIDIA devicenum=0

time(us): 4,802,950

51: data copyin reached 1 times

device time(us): total=22,768 max=22,768 min=22,768 avg=22,768

57: kernel launched 1000 times

grid: [4094] block: [256]

device time(us): total=2,611,387 max=2,817 min=2,593 avg=2,611

elapsed time(us): total=2,620,044 max=2,900 min=2,601 avg=2,620

57: reduction kernel launched 1000 times

grid: [1] block: [256]

device time(us): total=18,083 max=842 min=16 avg=18

elapsed time(us): total=27,731 max=852 min=25 avg=27

69: kernel launched 1000 times

grid: [4094] block: [256]

device time(us): total=2,130,162 max=2,599 min=2,112 avg=2,130

elapsed time(us): total=2,139,919 max=2,712 min=2,112 avg=2,130

83: data copyout reached 1 times

device time(us): total=20,550 max=20,550 min=20,550 avg=20,550

0.23 seconds


0.24 seconds

Execution Time (lower is better)



CPU: Intel i7-3930K
6 Cores @ 3.20GHz

GPU: NVIDIA Tesla K20



Further speedups



- OpenACC gives us more detailed control over parallelization
 - Via **gang**, **worker**, and **vector** clauses
- By understanding more about the specific GPU on which you're running, using these clauses may allow better performance.
- By understanding bottlenecks in the code via profiling, we can reorganize the code for even better performance
- More on this in the **Optimizing OpenACC** session from GTC2013

Communication & IO with OpenACC

Calling MPI with OpenACC (Standard MPI)



```
!$acc data copy(A)  
!$acc parallel loop
```

```
do i=1,N
```

```
...
```

```
enddo
```

```
!$acc end parallel loop
```

```
call neighbor_exchange(A)
```

```
!$acc parallel loop
```

```
do i=1,N
```

```
...
```

```
enddo
```

```
!$acc end parallel loop
```

```
!$acc end data
```

Array “A” resides in GPU memory.

Routine contains MPI and requires “A.”

Array “A” returns to CPU here.

OpenACC update Directive



Programmer specifies an array (or partial array) that should be refreshed within a data region.

```
do_something_on_device()
```

```
!$acc update host(a) ◀
```

Copy “a” from GPU to CPU

```
do_something_on_host()
```

```
!$acc update device(a) ◀
```

Copy “a” from CPU to GPU

The programmer may choose to specify only part of the array to update.

Calling MPI with OpenACC (Standard MPI)



```
!$acc data copy(A)
!$acc parallel loop
do i=1,N
  ...
enddo
!$acc end parallel loop
!$acc update host(A)
call neighbor_exchange(A)
!$acc update device(A)
!$acc parallel loop
do i=1,N
  ...
enddo
!$acc end parallel loop
!$acc end data
```



Copy “A” to CPU for MPI.



Return “A” after MPI to GPU.

OpenACC host_data Directive



Programmer specifies that host arrays should be used within this section, unless specified with **use_device**. This is useful when calling libraries that expect GPU pointers.

```
!$acc host_data use_device(a)
call MPI_Sendrecv(a,...)
!$acc end host_data
```

Pass the device copy of
“a” to subroutine.

```
#pragma host_data use_device(a)
{
  cublasDgemm(..., a, ...) ;
}
```

Pass the device copy of
“a” to function.

This directive allows interoperability with a variety of other technologies, CUDA, accelerated libraries, OpenGL, etc.

Calling MPI with OpenACC (GPU-aware MPI)



```
!$acc data copy(A)
!$acc parallel loop
do i=1,N
  ...
enddo
!$acc end parallel loop
!$acc host_data use_device(A)
call neighbor_exchange(A)
!$acc end host_data
!$acc parallel loop
do i=1,N
  ...
enddo
!$acc end parallel loop
!$acc end data
```

Pass device “A” directly
to a GPU-aware MPI
library called in
neighbor_exchange.

*More information about GPU-aware MPI libraries is available in other sessions, please see your agenda.

Thank you

