

Merge-like Functions and Load-Balancing Search

Sean Baxter

NVIDIA Research



Modern GPU (new)



<http://nvlabs.github.io/moderngpu> website

<https://www.github.com/nvlabs/moderngpu> code
fork me!

sbaxter@nvidia.com email

@moderngpu twitter

<http://www.moderngpu.com/stanford.pptx> slides

Modern GPU (2)



- <http://nvlabs.github.io/moderngpu>
- **New library and ebook for CUDA programming**
 - Bulk Insert/Bulk Remove
 - Merge
 - Mergesort
 - Segmented sort
 - Vectorized sorted search
 - Load-balancing search
 - IntervalMove (vectorized memcpy)
 - Relational joins
 - Multiset operations (eg set_difference)

Irregular problems



- **Show that GPU can handle irregular problems.**
- **Domain of GPU computing is more than embarrassingly-parallel problems.**
- **C++ STL array-processing functions are good starting point.**

Scheduling and decomposition



- **Challenge of irregular problems is scheduling or decomposition**
- **We map what work to which threads, and when do we run it?**
- **Scheduling not a real concern with sequential CPU programming. Central issue in massively-parallel programming.**

Serial Merge



```
template<typename T, typename Comp>
void CPUMerge(const T* a, int aCount, const T* b, int bCount,
              T* dest, Comp comp) {

    int count = aCount + bCount;
    int ai = 0, bi = 0;
    for(int i = 0; i < count; ++i) {
        bool p;
        if(bi >= bCount) p = true;
        else if(ai >= aCount) p = false;
        else p = !comp(b[bi], a[ai]); // a[ai] <= b[bi]

        dest[i] = p ? a[ai++] : b[bi++];
    }
}
```

Examine two keys and output one element per iteration. $O(n)$ work-efficiency.

Parallel Merge



- **PRAM-style merge**
 - Low-latency when number of processors is order N .
 - One item per thread. Communication free.
- **Two kernels:**
 1. **KernelA assigns one thread to each item in A.**
Thread i *lower-bound* binary searches into B for key $A[i]$.
Insert $A[i]$ into dest at $i + \text{lower_bound}(B, A[i])$.
 2. **KernelB assigns one thread to each item in B.**
Thread i *upper-bound* binary searches into A for key $B[i]$.
Insert $B[i]$ into dest at $i + \text{upper_bound}(A, B[i])$.

Parallel Merge (2)



```
template<int NT, typename T, typename Comp>
__global__ void ParallelMergeA(const T* a_global, int aCount,
    const T* b_global, int bCount, T* dest_global, Comp comp) {

    int gid = threadIdx.x + NT * blockIdx.x;
    if(gid < aCount) {
        T aKey = a_global[gid];
        int lb = BinarySearch<MgpuBoundsLower>(b_global, bCount,
            aKey, comp);
        dest_global[gid + lb] = aKey;
    }
}
```

ParallelMergeB is symmetric. Uses upper-bound for B into A search.

Parallel Merge (3)



- **Parallel version is highly concurrent but very inefficient.**
 - Serial code is $O(n)$
 - Parallel code is $O(n \log n)$
- **Parallel code doesn't resemble sequential code at all**
 - Hard to extend to other merge-like operations.
- **Parallel code tries to solve two problems at once:**
 1. Decomposition/scheduling work to parallel processors.
 2. Merge-specific logic

Two-phase decomposition



- **Design implementation in two phases:**

- 1. PARTITIONING PHASE**

Maps work onto each CTA or thread.

Has adjustable *grain size* parameter (VT).

Implemented with one binary search per CTA or thread.


- 2. WORK LOGIC**

Executes code specific for solving problem.

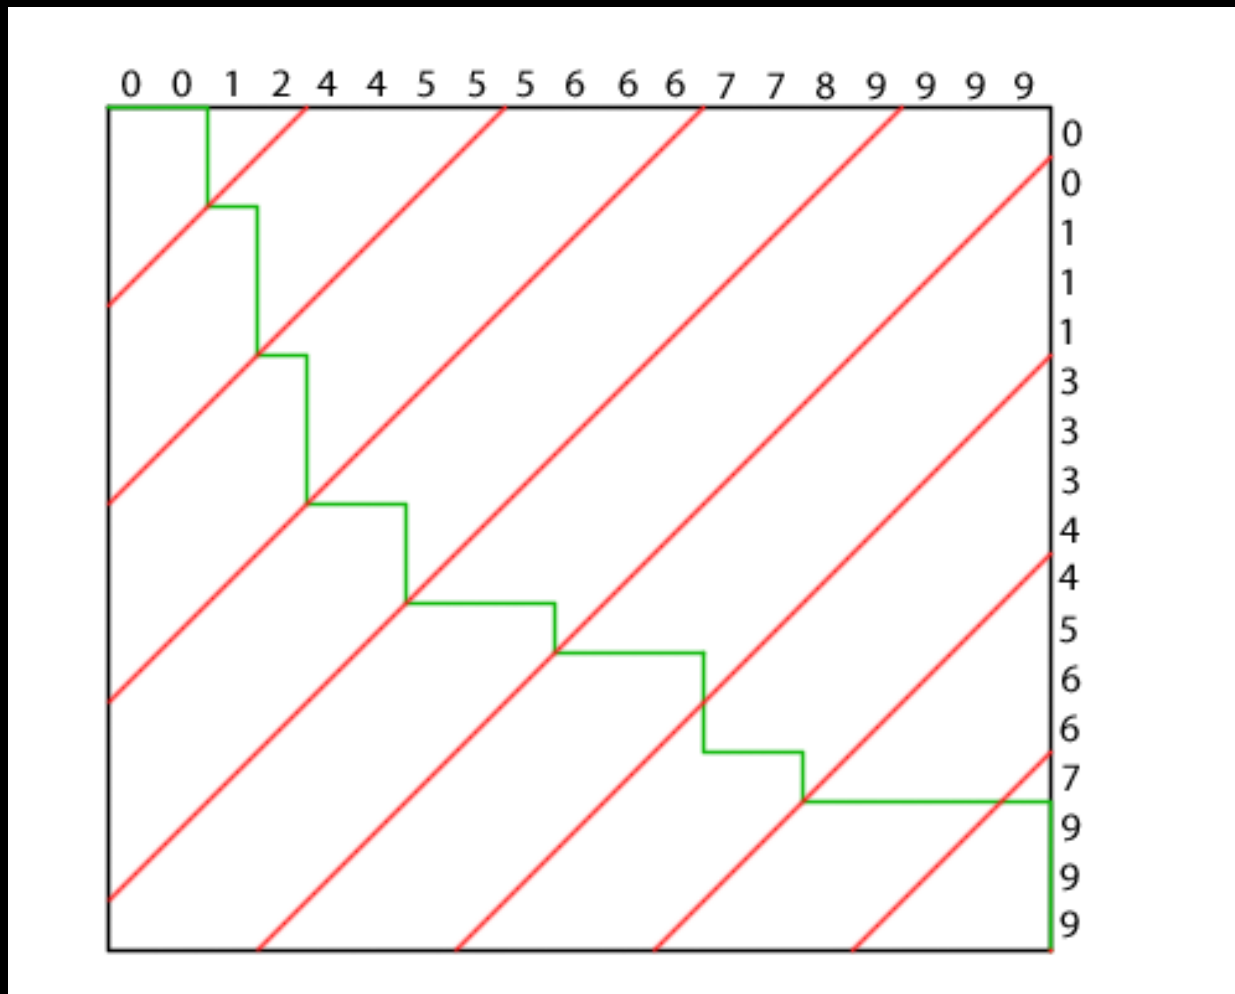
This should resemble CPU sequential implementation.

More work-efficient and more extensible.

Merge Path



Merge Path (2)



Binary Search



```
template<typename T, typename It, typename Comp>
    int BinarySearch(It data, int count, T key, Comp comp) {
    int begin = 0;
    int end = count;
    while(begin < end) {
        int mid = (begin + end)>> 1;
        bool pred = comp(key2, key);    // key <= key2.
        if(pred) begin = mid + 1;
        else end = mid;
    }
    return begin;
}
```

Standard binary search with support for C++ comparators. Searches for a key in one array.

Merge Path (3)

```
template<typename T, typename It1, typename It2, typename Comp>
int MergePath(It1 a, int aCount, It2 b, int bCount, int diag,
              Comp comp) {
    int begin = max(0, diag - bCount);
    int end = min(diag, aCount);

    while(begin < end) {
        int mid = (begin + end) >> 1;
        bool pred = !comp(b[diag - 1 - mid], a[mid]);
        if(pred) begin = mid + 1;
        else end = mid;
    }
    return begin;
}
```

Merge Path binary search with support for C++ comparators. Simultaneously search two arrays by using **constraint $a_i + b_i = \text{diag}$ to make problem one dimensional.**

Parallel Merge (4)



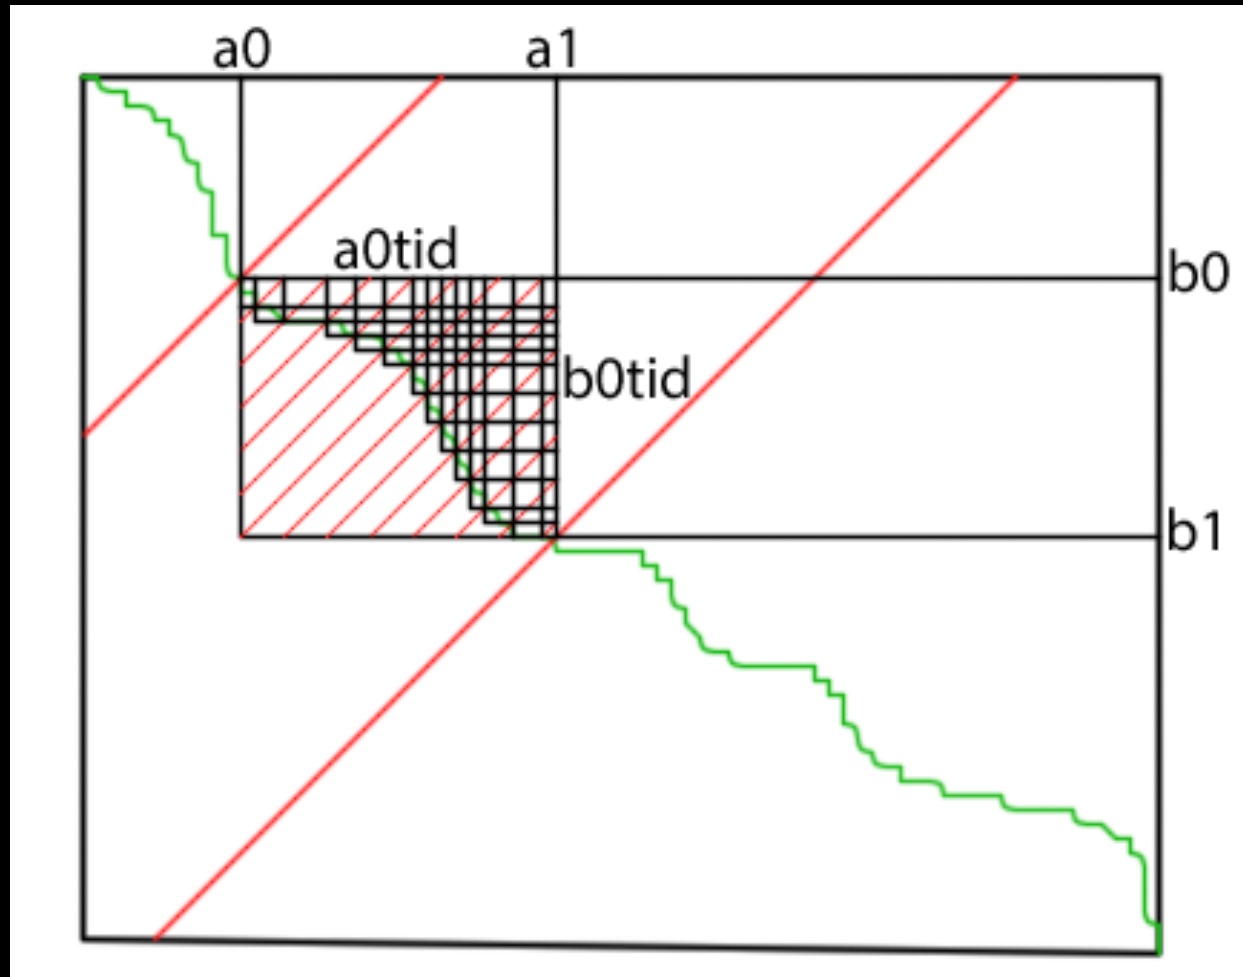
1. Partitioning phase:

1. Use MergePath search to partition input arrays into CTA-sized intervals.
2. In CTA, load interval from A and interval from B.
3. Run MergePath in CTA into shared memory to partition work per thread.

2. Work phase:

Use SerialMerge to merge VT elements per thread without communication.

Parallel Merge (5)



Serial Merge


```
#pragma unroll
for(int i = 0; i < Count; ++i) {
    T x1 = keys[aBegin];
    T x2 = keys[bBegin];

    // If p is true, emit from A, otherwise emit from B.
    bool p;
    if(bBegin >= bEnd) p = true;
    else if(aBegin >= aEnd) p = false;
    else p = !comp(x2, x1);          // p = x1 <= x2

    // because of #pragma unroll, merged[i] is static indexing
    // so results is kept in RF, not smem!
    results[i] = p ? x1 : x2;
    if(p) ++aBegin;
    else ++bBegin;
}
```

- SerialMerge processes a *constant* number of elements per thread.
- Unroll loops and store to register.

MGPU Merge throughput



Parallel Merge (6)




- **SerialMerge is looks like CPUMerge.**
- **It's work-efficient.**
- **It's easy to extend.**
- **Two-phase design presents opportunity for tuning:**
 - **Search tuning space of grain size parameter VT.**
 - **Increase VT to do more serial work per search and improve work-efficiency (amortizes partitioning cost).**
 - **Decrease VT to increase occupancy and achieve better execution efficiency.**

Tuning



Mergesort Throughputs
Parameter Tunings



Tuning (2)



- **Kepler has a wider SM-needs more parallelism to keep execution units fed.**
- **Fermi has more shared memory per execution unit to hide instruction latency. Choose larger grain size beneficial.**

	<u>GTX 480</u> <u>(Fermi)</u>	<u>GTX Titan</u> <u>(Kepler)</u>
32-bit int	128x23	256x11
64-bit int	128x11	256x5

Merge-like Functions

- Merge (and mergesort)
- Segmented sort
- Sorted searches – sorted needles in sorted haystack
 - lower_bound
 - upper_bound
 - equal_range
 - counts
- Sets
 - set_intersection
 - set_union
 - set_difference
 - set_symmetric_difference
- Joins (built on sorted searches)
 - Inner, left, right, full
 - Semi-join, anti-join

Vectorized sorted search

- Unusual new primitive.
- Searches array of sorted keys into array of sorted haystack.
- Simple usage: return lower-bound (or upper-bound) of A into B.
- Power usage:
 - Lower-bound of A into B.
 - Upper-bound of B into A.
 - Flags for all matches of A into B.
 - Flags for all matches of B into A.
 - All this with a single pass over data!

Vectorized sorted search (2)

```
// Return lower-bound of A into B.
template<typename T, typename Comp>
void CPUSortedSearch(const T* a, int aCount, const T* b,
    int bCount, int* indices, Comp comp) {

    int aIndex = 0, bIndex = 0;
    while(aIndex < aCount) {
        bool p;
        if(bIndex >= bCount) p = true;
        else p = !comp(b[bIndex], a[aIndex]);

        if(p) indices[aIndex++] = bIndex;
        else ++bIndex;
    }
}
```

Looks just like CPU Merge!

Vectorized sorted search (3)

- Parallelize sorted search just like we parallelized merge.
- Use same MergePath code for partitioning problem.
- Use different serial work logic code to specialize for this problem.
 - Two-phase decomposition promotes code reuse.
- Coarse-grained partitioning maps a fixed amount of *needles plus haystack* into each CTA.
- Fine-grained partitioning maps a fixed amount of *needles plus haystack* into each thread.

SerialSearch

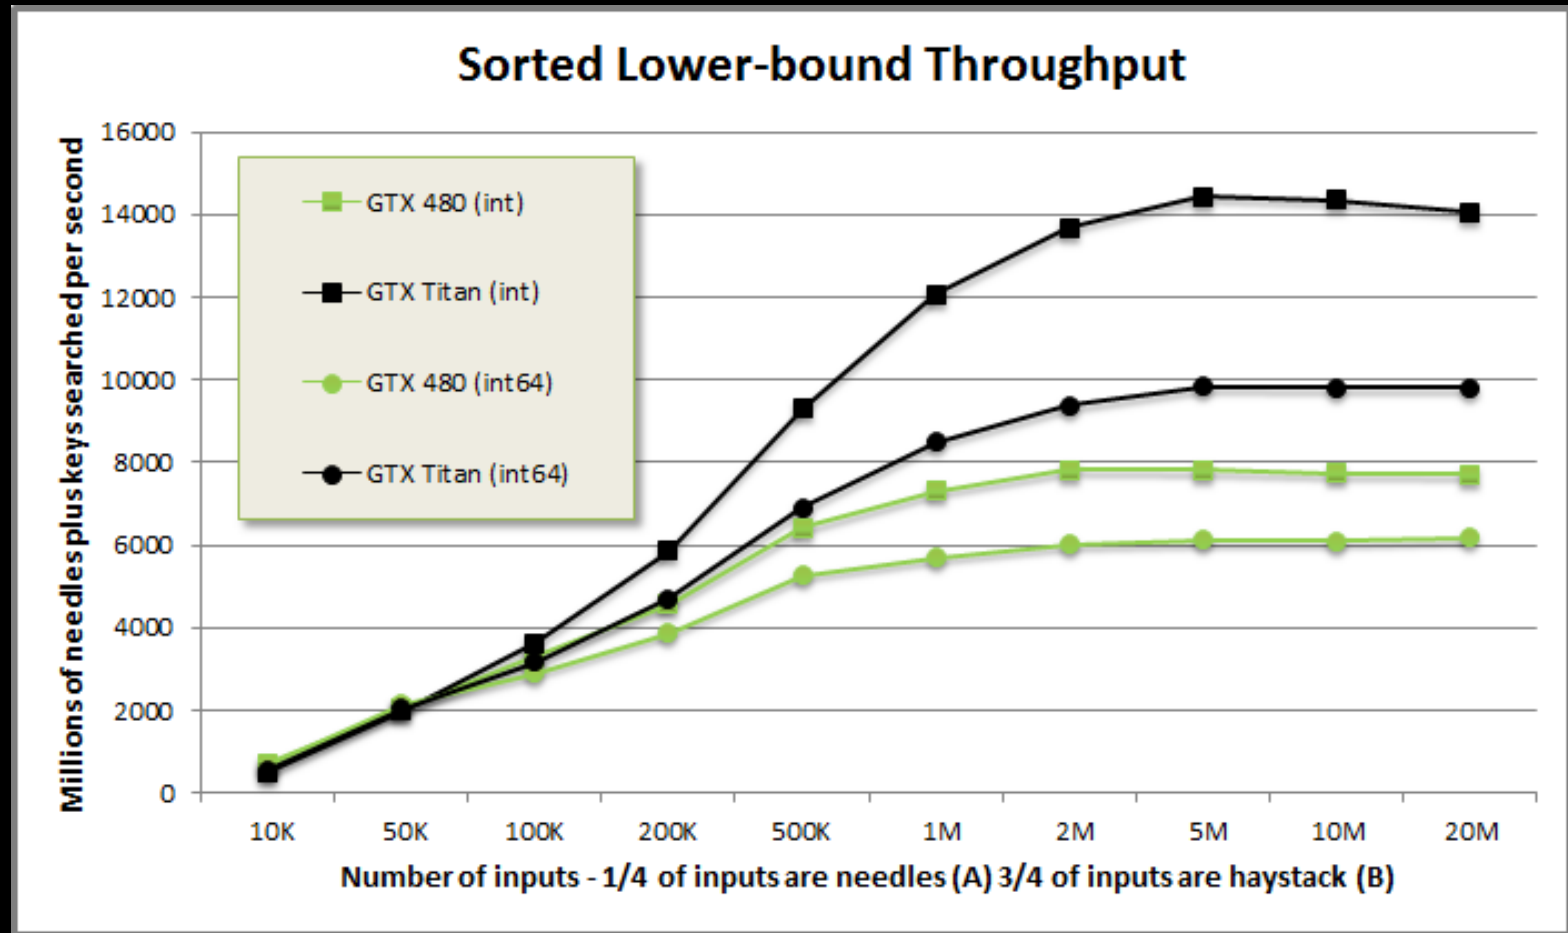


```
template<int VT, typename T, typename Comp>
MGPU_DEVICE int DeviceSerialSearch(const T* keys_shared,
    int aBegin, int aEnd, int bBegin, int bEnd, int bOffset,
    int* indices, Comp comp) {


    int decisions = 0;
    #pragma unroll
    for(int i = 0; i < VT; ++i) {
        bool p;
        if(aBegin >= aEnd) p = false;
        else if(bBegin >= bEnd) p = true;
        else p = !comp(keys_shared[bBegin], keys_shared[aBegin]);

        if(p) {
            decisions |= 1<< i; // Set bit to indicate result
            indices[i] = bOffset + bBegin;
            ++aBegin;
        } else ++bBegin;
    }
    return decisions;
}
```

Vectorized sorted search throughput



Vectorized sorted search throughput



Load-balancing search



- **Load-balancing search is a special decomposition**
- **... Or a change of coordinates**
- **... Or a kind of inverse of prefix sum**
- **N objects**
 - **Each object generates variable outputs**
 - **Match each output with the generating object.**

Load-balancing search (2)



Work-item counts:

0:	1	2	4	0	4	4	3	3	2	4
10:	0	0	1	2	1	1	0	2	2	1
20:	1	4	2	3	2	2	1	1	3	0
30:	2	1	1	3	4	2	2	4	0	4

Exc-scan of counts:

0:	0	1	3	7	7	11	15	18	21	23
10:	27	27	27	28	30	31	32	32	34	36
20:	37	38	42	44	47	49	51	52	53	56
30:	56	58	59	60	63	67	69	71	75	75

Total work-items: 79

Load-balancing search (3)



Load-balancing search:

0:	0	1	1	2	2	2	2	4	4	4
10:	4	5	5	5	5	6	6	6	7	7
20:	7	8	8	9	9	9	9	12	13	13
30:	14	15	17	17	18	18	19	20	21	21
40:	21	21	22	22	23	23	23	24	24	25
50:	25	26	27	28	28	28	30	30	31	32
60:	33	33	33	34	34	34	34	35	35	36
70:	36	37	37	37	37	39	39	39	39	

Work-item rank:

0:	0	0	1	0	1	2	3	0	1	2
10:	3	0	1	2	3	0	1	2	0	1
20:	2	0	1	0	1	2	3	0	0	1
30:	0	0	0	1	0	1	0	0	0	1
40:	2	3	0	1	0	1	2	0	1	0
50:	1	0	0	0	1	2	0	1	0	0
60:	0	1	2	0	1	2	3	0	1	0
70:	1	0	1	2	3	0	1	2	3	

Load-balancing search (4)

- Each output is paired with its generating object
- A rank for the work-item within the generating object is also returned.
- LBS is computed as the *upper-bound of the natural numbers into the scan of the work-item counts minus 1*.
- Use vectorized sorted search (upper-bound) pattern with some optimizations.
- Use ordinary Merge Path search for coarse-grained partitioning.

Load-balancing search (5)


```
template<int VT>
MGPU_DEVICE void DeviceSerialLoadBalanceSearch(
    const int* b_shared, int aBegin, int aEnd, int bFirst,
    int bBegin, int bEnd, int* a_shared) {

    int bKey = b_shared[bBegin];

    #pragma unroll
    for(int i = 0; i < VT; ++i) {
        bool p = (aBegin < aEnd) &&
            ((bBegin >= bEnd) || (aBegin < bKey));

        if(p)          // Advance A (the needle).
            a_shared[aBegin++] = bFirst + bBegin;
        else            // Advance B (the haystack).
            bKey = b_shared[++bBegin];
    }
}
```

Load-balancing search throughput



Load-balancing search (6)



- Run CTALoadBalance inside kernel as boilerplate.
- This transforms the problem so that its dependencies are explicit.
- Great for implementing functions that *expand* or *contract* data:
 - IntervalMove (vectorized cudaMemcpy)
 - Relational joins
 - Sparse matrix * matrix (expand partials)

IntervalMove




- **Many small cudaMemcpy calls:**
 - Too slow.
- **One kernel:**
 - Batch up many cudaMemcpy's by gather (source), scatter (destination), and size parameters, and fire off in a single launch.
- **Scan interval counts and use *load-balancing search* to map output elements into input intervals.**
- **Launch enough KernelIntervalMove threads to cover both the outputs *and* inputs.**

IntervalMove (2)

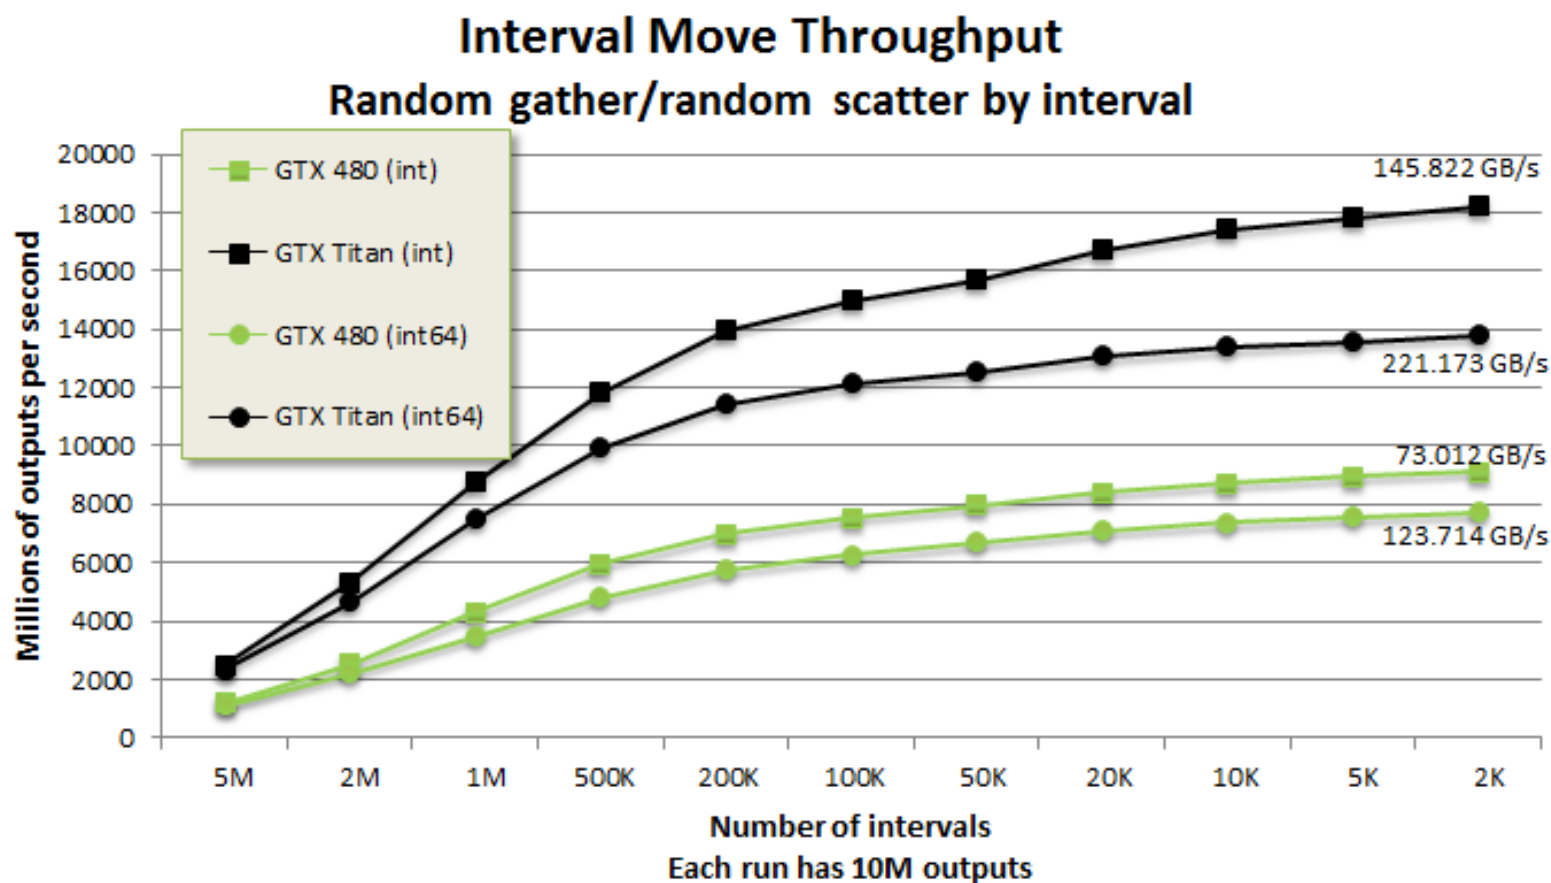


- **CTALoadBalance inside KernelIntervalMove maps up to VT outputs per thread into the interval that generated it.**
- **The caller infers the rank of each output within the interval to know which element within the interval to copy.**
- **Because of Merge Path's exact partitioning, each thread copies no more than VT elements, no matter the distribution of intervals.**

IntervalMove throughput



IntervalMove throughput



Relational Joins



- MGPU implements full-outer join on GPU with exact load-balancing.

<u>Row</u>	<u>A index</u>	<u>A key</u>	<u>B key</u>	<u>B index</u>	<u>Join type</u>
0	0	A ⁰	A ⁰	0	inner
1	0	A ⁰	A ¹	1	inner
2	1	A ¹	A ⁰	0	inner
3	1	A ¹	A ¹	1	inner
4	2	B ⁰	B ⁰	2	inner
5	2	B ⁰	B ¹	3	inner
6	2	B ⁰	B ²	4	inner
7	3	E ⁰	---	-1	left
8	4	E ¹	---	-1	left
9	5	E ²	---	-1	left
10	6	E ³	---	-1	left
11	7	F ⁰	F ⁰	7	inner
12	8	F ¹	F ⁰	7	inner
13	9	G ⁰	G ⁰	8	inner
14	9	G ⁰	G ¹	9	inner
15	10	H ⁰	H ⁰	10	inner
16	11	H ¹	H ⁰	10	inner
17	12	J ⁰	---	-1	left
18	13	J ¹	---	-1	left
19	14	M ⁰	---	-1	left
20	15	M ¹	---	-1	left
21	-1	---	C ⁰	5	right
22	-1	---	C ¹	6	right
23	-1	---	I ⁰	11	right
24	-1	---	L ⁰	12	right
25	-1	---	L ¹	13	right

Relational Join (2)




- We join two sorted tables (sort-merge join).
- Equal keys in A and B are expanded with Cartesian product.
- Keys in A not found in B are emitted with left-join (null B key).
- Keys in B not found in A are emitted with right-join (null A key).

Relational Join (3)



- **Vectorized lower-bound and upper-bound search finds set of matches for each row in A in the B table.**
- **Difference of upper- and lower-bound iterators are match *counts*.**
- **Scan match counts.**
- **Use load-balancing search**

Relational join throughputs




Segmented sort



- **MGPU Merge implements MGPU Mergesort.**
- **Extend Merge Path decomposition to segments for MGPU Segmented Sort.**
- **Simultaneously sort many variable-length arrays.**
- **Without compromising concurrency, code improves over $O(n \log n)$ complexity by detecting segment structure and early-exiting.**

Segmented sort throughput



Wrap-up



- **Separate partitioning from work logic.**
- **Expose grain size in partitioning for tuning.**
- **Write work-efficient, sequential, communication-free code to implement function.**
- **Exploit *exact partitioning***
 - **Unroll all loops.**
 - **Read from shared memory and store to register.**
 - **Reduce smem consumption and boost occupancy.**

Wrap-up (2)



- **Use load-balancing search to accommodate expansion of contraction of data from per-object counts.**
- **Use CTALoadBalance to embed load-balancing search (itself a two-phase vectorized upper-bound search) as boilerplate in kernels.**
- **Use IntervalMove (or IntervalGather or IntervalScatter) to queue up and launch many cudaMemcpy operations at once.**

Questions?



<http://nvlabs.github.io/moderngpu> website

<https://www.github.com/nvlabs/moderngpu> code
fork me!

sbaxter@nvidia.com email

@moderngpu twitter

<http://www.moderngpu.com/stanford.pptx> slides